

FAST GRAPH CLASSIFIER WITH OPTICAL RANDOM FEATURES

Hashem Ghanem Nicolas Keriven Nicolas Tremblay

CNRS, GIPSA-lab, FR-38402 Saint Martin d’Heres Cedex, France

ABSTRACT

Graph Classification is the problem of mapping graphs to the classes they belong to. The graphlet kernel is a classical method in graph classification. However, it suffers from high computation cost due to the isomorphism test it includes. We propose a generic algorithm that mainly replaces the isomorphism test with a user-defined mapping. This algorithm computes a feature vector for each graph. More importantly, we incorporate *optical processing units (OPUs)* technology by choosing the optical random features map. We show that OPUs complete such mapping in *light-speed*. Besides, We prove theoretically that using the OPU is effective in graph classification. Finally, we support this with the necessary experiments.

Index Terms— Optical random features, Graph kernels

1. INTRODUCTION

Graph structures are used to model a set of objects and their interactions. In biology for instance, proteins can be modeled as graphs where amino acids are nodes and the chemical links between them are edges [1]. In social networks, each post thread is modeled with a graph whose nodes are users, where an edge between two nodes exists if one replies to the other’s comment on that thread [2].

Having *a priori* known graph classes, the graph classification task is to predict the class of any given graph. In biology, proteins are to be classified to enzymes and to non-enzymes [1], as in the D&D dataset. In social networks, one task is to discriminate between discussion-based and question/answer-based threads [3], as in Reddit-Binary dataset.

In addition to the graph structure, nodes and edges may have extra features that can be used in classifying graphs. It has been shown that node features are important to have high performance [4]. Here, we focus on the slightly harder case where one has only access to the graph structure.

Moreover, we place ourselves in the context of *supervised learning*, where we have access to a set of pre-labeled graphs ($\mathcal{X} = \{\mathcal{G}_1, \dots, \mathcal{G}_n\}$, $\mathcal{Y} = \{y_1, \dots, y_n\}$). Each graph \mathcal{G}_i is a *a priori* known to belong to the class with label y_i . The graph classification problem is: given this prior information, design an algorithm that, given in input a new graph, outputs the label of the class to which it belongs.

Related work Structure-based graph classification has been tackled using many algorithms. Frequent subgraphs based algorithms, gSpan method [5] for instance, perform a prohibitive-cost analysis on the graph dataset \mathcal{X} to catch the frequent and discriminative subgraphs, then use them as features. Graph kernel-based algorithms [6] compute fixed representation for graphs by defining the kernel as a similarity function between graphs. One example is *the graphlet kernel*, which is an efficient method in the literature but is prone to high computation cost [7]. Graph convolutional networks (GCNs) are used to solve graph classification too. Recently, a particular model called GIN (Graph Isomorphism Network) was developed and provided high performance classification [8].

Contribution: Random features are an computationally efficient method to approximate kernel functions which satisfy some conditions [9]. Optical Processing Units (OPUs) is a recently developed hardware that computes optical random features in light speed. On the other hand, the graphlet kernel represents a graph by how many times graphs of a smaller size occur in it. This kernel includes the isomorphism test in the counting process which makes it computationally expensive. Can we leverage OPUs computations to reduce the computational complexity of a combinatorial problem like the graphlet kernel ? We show, empirically and theoretically, that yes we can do that to obtain fast and efficient graph classifier.

2. BACKGROUND

2.1. The graphlet kernel

First, we present the concepts necessary to define the graphlet kernel. We represent a graph of size v by the adjacency matrix $\mathbf{A} \in \mathbb{R}^{v \times v}$, which has $a_{i,j} = 1$ if there is an edge between nodes $\{i, j\}$ and 0 otherwise. Two graphs are said to be isomorphic ($\mathcal{G} \cong \mathcal{G}'$) iff we can permute the nodes’ labels of one such that their adjacency matrices are equal [10].

We denote by $\mathfrak{H} = \{\mathcal{H}_1, \dots, \mathcal{H}_{N_k}\}$ the set of all non-isomorphic graphs, also called graphlets, of size k . We define the matching function $\varphi_k^{match}(\mathcal{F}) = [1_{(\mathcal{F} \cong \mathcal{H}_i)}]_{i=1}^{N_k} \in \{0, 1\}^{N_k}$, where 1_{Ω} is the indicator function. In words, $\varphi_k^{match}(\mathcal{F})$ is a Boolean vector of dimension N_k which has a 1 in the coordinate i if $\mathcal{F} \cong \mathcal{H}_i$, and 0 otherwise. Let $\mathfrak{F}_{\mathcal{G}} = \{\mathcal{F}_1, \mathcal{F}_2, \dots\}$ be the set of subgraphs induced by all size- k subsets of nodes in a graph \mathcal{G} . We assign each graph

with the following representation vector, called k -spectrum:

$$\mathbf{f}_G = \frac{1}{|\mathcal{F}_G|} \sum_{\mathcal{F} \in \mathcal{F}_G} \varphi_k^{\text{match}}(\mathcal{F}) \in \mathbb{R}^{N_k} \quad (1)$$

For two graphs $\mathcal{G}, \mathcal{G}'$, the graphlet kernel is defined by the inner product $\mathbf{f}_G^T \mathbf{f}_{G'}$. The graphlet kernel performs well especially with a sufficiently large value of k [7]. However, in each graph \mathcal{G} of size v , the computation cost to compute \mathbf{f}_G is $C_{gk} = \mathcal{O}\left(\binom{v}{k} N_k C_k^{\cong}\right)$, where C_k^{\cong} is the cost of the isomorphism test between two graphs of size k . This cost is expensive due to: i/ $\binom{v}{k}$ explodes as v or k increase, ii/ N_k is exponential in k , iii/ yet there is no known method to test isomorphism in polynomial time of k [11].

Usually, uniform graphlet sampling is used to accelerate the graphlet kernel [7], where sampling a size- k subgraph means: first we uniformly at random choose k nodes from the graph, then we sample the subgraph induced by those nodes. This yields that each sample follow a uniform distribution over \mathcal{F}_G .

Knowing this, the k -spectrum \mathbf{f}_G can be interpreted as follows: if one samples a subgraph from \mathcal{G} , then one has a probability $(\mathbf{f}_G)_i$ of obtaining \mathcal{H}_i , i.e.: $\mathbf{f}_G = \mathbb{E}_{F \sim \text{unif}} \varphi_k^{\text{match}}(F)$. It is thus natural to approach \mathbf{f}_G with a sample average, where by sampling s subgraphs of size k to form the collection $\tilde{\mathcal{F}}_G = \{F_1, \dots, F_s\}$, the estimator:

$$\hat{\mathbf{f}}_G = \frac{1}{s} \sum_{F \in \tilde{\mathcal{F}}_G} \varphi_k^{\text{match}}(F). \quad (2)$$

verifies by the law of large numbers that $\hat{\mathbf{f}}_G \xrightarrow{s \rightarrow \infty} \mathbf{f}_G$ with probability 1.

The computation cost per graph of the graphlet kernel with graph sampling is $C_{gk+gs} = \mathcal{O}(s C_S N_k C_k^{\cong})$, where C_S is the cost of sampling one subgraph. Although the term $\binom{v}{k}$ doesn't exist in this cost, still for a specific certainty in estimating \mathbf{f}_G , the required number of samples s must be proportional to N_k [7]. So this version is still expensive especially when k is large.

In fact, there exist many different sampling techniques S_k , each follows a specific random process in sampling the k nodes from a graph \mathcal{G} . As a result, subgraphs sampled with a technique different from the uniform sampling will have a different histogram \mathbf{f}_{G, S_k} than the one defined by \mathbf{f}_G [12]. Such sampling technique is the random walk (RW) sampler. Unlike uniform sampling, RW tends to sample connected subgraphs, which are more informative about the graph structure.

3. METHOD

3.1. Proposed algorithm

We propose to replace φ_k^{match} with a user-defined map φ , and to replace the uniform sampler with a user-chosen one like

Algorithm 1: GSA- φ generic algorithm

Input: labeled graph dataset $\mathcal{X} = (\mathcal{G}_i, y_i)_{i=1, \dots, n}$
1 Tools Graphlet sampler S_k , a function φ , linear classifier (ex. SVM)
2 Hyperparameters k : graphlet size, s : number of graphlet samples per graph, m : number of random features
Output: Trained model to classify graphs
3 Algorithm
4 Random initialization of the SVM weights
5 for \mathcal{G}_i in \mathcal{X} **do**
6 $\mathbf{z}_i = \mathbf{0}$ (null vector of size m)
7 **for** $j = 1 : s$ **do**
8 $F_{i,j} \leftarrow S_k(\mathcal{G}_i)$
9 $\mathbf{z}_i \leftarrow \mathbf{z}_i + \frac{1}{s} \varphi(F_{i,j})$
10 $\mathcal{D}_\varphi \leftarrow (\mathbf{z}_i, y_i)_{i=1, \dots, n}$
11 Train the linear classifier on the new vector-valued dataset \mathcal{D}_φ

the RW sampler. This makes it possible to deploy the OPU technology with this algorithm. We refer to this framework as *graphlet sampling and Averaging GSA* – φ , and its cost per graph for a specific φ is $C_{GSA-\varphi} = \mathcal{O}(s C_S C_\varphi)$, where C_φ is the cost of applying φ on one graphlet.

The function φ here maps each subgraph to an m -dimensional space \mathbb{R}^m . Note that choosing $\varphi = \varphi_k^{\text{match}}$ and S_k as the uniform sampler, GSA- φ_k^{match} turns out to be the graphlet kernel with graphlet sampling. We see next that choosing φ as random maps is both fast and efficient.

3.2. Efficiency of kernel random features with GSA – φ

A kernel κ associated to a random features (RF) decomposition is a positive definite function of two inputs that can be decomposed as follows [13]:

$$\kappa(\mathbf{x}, \mathbf{x}') = \mathbb{E}_{\mathbf{w} \sim p} [\xi_{\mathbf{w}}(\mathbf{x}) \xi_{\mathbf{w}}(\mathbf{x}')] \quad (3)$$

where \mathbb{E} stands for the expectation, p is a probability distribution, and ξ is a function parameterized by \mathbf{w} . To approximate such kernels, we can empirically average m realizations of $\xi_{\mathbf{w}}(\mathbf{x}) \xi_{\mathbf{w}}(\mathbf{x}')$. To do that, we define RF maps:

$$\varphi(\mathbf{x}) = \frac{1}{\sqrt{m}} (\xi_{\mathbf{w}_j}(\mathbf{x}))_{j=1}^m \in \mathbb{C}^m \quad (4)$$

where m is called here the number of features and the frequencies \mathbf{w}_j are drawn identically and independently (iid) from p . Then, $\kappa(\mathbf{x}, \mathbf{x}') \approx \varphi(\mathbf{x})^T \varphi(\mathbf{x}')$.

For each kernel as in (3), and for any two graphs $\mathcal{G}, \mathcal{G}'$, we define the MMD (Maximum Mean Discrepancy) metric between the two corresponding distributions $\mathbf{f}_{G, S_k}, \mathbf{f}_{G', S_k}$ by:

$$\text{MMD}^2(\mathbf{f}_{G, S_k}, \mathbf{f}_{G', S_k}) = \mathbb{E}_{\mathbf{w}} \left(\left| \mathbb{E}_{S_k(\mathcal{G})} \xi_{\mathbf{w}}(F) - \mathbb{E}_{S_k(\mathcal{G}')} \xi_{\mathbf{w}}(F') \right|^2 \right)$$

The main property of the MMD is that, for so-called *characteristic kernels*, it is a true metric on distributions, *i.e.* $MMD(\mathcal{P}, \mathcal{Q}) = 0 \Leftrightarrow \mathcal{P} = \mathcal{Q}$. In addition, most usual kernels, like the Gaussian kernel, are characteristic [14].

The next theorem shows the efficiency of using RF maps as in (4) with our algorithm. It states that the Euclidean distance between the representation vectors $\mathbf{z}_G, \mathbf{z}_{G'}$ converges to the MMD metric between their distributions $\mathbf{f}_{G, S_k}, \mathbf{f}_{G', S_k}$.

Theorem 1. *Let \mathcal{G} and \mathcal{G}' be two graphs, $\mathfrak{F}_G = \{F_i\}_{i=1}^s$ (resp. $\mathfrak{F}_{G'} = \{F'_i\}_{i=1}^s$) be iid size- k graphlet samples drawn from $S_k(\mathcal{G})$ (resp. $S_k(\mathcal{G}')$). Assume a random feature map as in (4). Assume that $|\xi_{\mathbf{w}}(F)| \leq 1$ for any \mathbf{w}, F . We have for all $\delta > 0$ and with probability at least $1 - \delta$:*

$$\left| \|\varphi(\mathfrak{F}_G) - \varphi(\mathfrak{F}_{G'})\|^2 - MMD(\mathbf{f}_{G, S_k}, \mathbf{f}_{G', S_k})^2 \right| \leq \frac{4\sqrt{\log(6/\delta)}}{\sqrt{m}} + \frac{8\left(1 + \sqrt{2\log(3/\delta)}\right)}{\sqrt{s}}$$

Proof. see Appendix. \square

3.3. Considered choices of φ_{RF}

Gaussian maps φ_{Gs} applied on the adjacency matrix: φ_{Gs} is the RF map of the Gaussian kernel [13]. We for each subgraph \mathcal{F} take its vectorized adjacency matrix $\mathbf{a}_{\mathcal{F}} = \text{flatten}(\mathbf{A}_{\mathcal{F}})$ as input. Then:

$$\varphi_{Gs}(\mathcal{F}) = \frac{1}{\sqrt{m}} \left(\sqrt{2} \cos(\mathbf{w}_j^T \mathbf{a}_{\mathcal{F}} + b_j) \right)_{j=1}^m \in \mathbb{R}^m \quad (5)$$

where the frequencies $\mathbf{w}_j \in \mathbb{R}^{k^2}$ are drawn from a Gaussian distribution with the inverse variance of the original kernel.

Gaussian maps φ_{Gs+Eig} applied on the sorted eigenvalues of the adjacency matrix: instead of passing the vectorized adjacency matrix as input, we pass the vector of its sorted eigenvalues $\boldsymbol{\lambda} \in \mathbb{R}^k$. The motive proposing φ_{Gs+Eig} is that it respects the isomorphism test since: $\text{if } \lambda(\mathbf{A}) = \lambda(\mathbf{PAP}^T)$ for any permutation matrix \mathbf{P} , $\text{if } F \cong F' \Rightarrow \exists \mathbf{P}, \mathbf{A}_F = \mathbf{PA}_{F'}\mathbf{P}^T$. Thus, $F \cong F' \Rightarrow \varphi_{Gs+Eig}(F) = \varphi_{Gs+Eig}(F')$. φ_{Gs+Eig} maps isomorphic subgraphs to the same point in \mathbb{R}^m .

Optical random feature maps φ_{OPU} : This corresponds to the fastest version of our algorithm. OPUs (Optical Processing Units) technology was developed to compute a specific random features mapping in *constant time* $\mathcal{O}(1)$ in both m and k using light scattering [15]. Having the random matrix \mathbf{W} , traditional random maps (ex. φ_{Gs}) need $\mathcal{O}(mk^2)$ cost to compute $\mathbf{W}\mathbf{x}$ as in (5). An OPU computes its associated map at the speed of light, this map is modeled as follows [15]:

$$\varphi_{OPU}(\mathbf{x}) = |\mathbf{W}\mathbf{x} + \mathbf{b}|^2; \quad \mathbf{W} \in \mathbb{R}^{m \times d}, \mathbf{b} \in \mathbb{R}^m, \mathbf{x} \in \mathbb{R}^d$$

Where \mathbf{b} is a random bias vector, d is the input space dimension, the amplitude function $|\cdot|$ is taken element wise, and

Graphlet kernel		$O\left(\binom{v}{k} N_k C_k^{\infty}\right)$
GSA- φ with:	φ_k^{match}	$O(C_s s N_k C_k^{\infty})$
	φ_{Gs}	$O(C_s s m k^2)$
	φ_{Gs+Eig}	$O(C_s s (m k + k^3))$
	φ_{OPU}	$O(C_s s)$

Table 1. Per-graph complexities of GSA- φ .

\mathbf{W} is a random iid complex matrix with Gaussian real and imaginary parts. The complexities of the different mappings φ examined in this work are summarized in Table 1.

4. EXPERIMENTS

4.1. Setup

With respect to the performance and the computation time, we compare different choices of map φ in $GSA - \varphi$. We benchmark the performance of $GSA - \varphi_{OPU}$ against GIN based graph convolutional network [8].

In all experiments except the last two, we use a synthetic dataset generated by a *Stochastic Block Model (SBM)* [16]. We generate 300 graphs, 240 for training and 60 for testing. Each graph has $v = 60$ nodes divided equally between six communities. Moreover, graphs are divided into two classes $\{0, 1\}$ based on the edges distribution considered. For each class we fix two values (p_{in}, p_{out}) which are the probabilities of generating an edge between any two nodes when they are in the same community and when they are in different ones, respectively. Besides, to prevent the classes from being easily discriminated by the average degree, the pairs $(p_{in,i}, p_{out,i})_{i=0,1}$ are chosen such that all nodes have a fixed expected average degree equal to 10. Having one degree of freedom left, we fix $p_{in,1} = 0.3$, and we vary $r = (p_{in,1}/p_{in,0})$ the inter-class similarity parameter: the closer r is to 1, the more similar both classes are, and thus the harder it is to discriminate them.

In addition, D&D [17] and Reddit-Binary [3] are labeled datasets of size $n = 1178$, $n = 2000$ respectively. In D&D, nodes have 7 features each, which are not used by our algorithms, *i.e.* we classify graphs based on their structure only. In what follows, unless otherwise indicated, we use uniform sampling and the adjacency matrix of subgraphs as input.

4.2. Choice of feature map φ

Comparison of random features: Fig 1(a) shows that $GSA - \varphi_{OPU}$ gives a better test accuracy with sufficiently large m than both $GSA - \varphi_{Gs+Eig}$ or $GSA - \varphi_{Gs}$, whose variance σ^2 was chosen empirically based on the validation accuracy. On the contrary, $GSA - \varphi_{Gs+Eig}$ performs best with small values of m , but increasing this number does not really improve the result and it is over-matched at high m . A possible justification is that the eigenvalues of the adjacency




Fig. 1. Comparing different φ maps in $GSA - \varphi$. (a) test accuracy when using RF maps with $k = 6$ while varying m . (b) Computation time as a function of k . If not specified: $r = 1.1$, $s = 2000$, $m = 5000$ and the Gaussian map variance $\sigma^2 = 0.01$.

matrix lose information about the subgraphs, even though respecting the isomorphism means that we are working with a smaller histogram and less random features are required.

Comparing $GSA - \varphi_{OPU}$ to $GSA - \varphi_k^{match}$: from Fig 2 we observe that with the same limited number of samples s , $GSA - \varphi_{OPU}$ with both (uniform sampling, $k = 6$) and (RW sampling, $k = 5$) clearly outperforms the graphlet kernel with graphlet sampling $GSA - \varphi_k^{match}$ with $k = 6$. We conclude that $GSA - \varphi_{OPU}$ is more adapted in this case than the traditional graphlet kernel.

Computational time: Fig 1(b) shows the computation time per subgraph of previous methods with respect to the subgraph size k . Other parameters are identically fixed for all methods. As expected, the execution time of $GSA - \varphi_k^{match}$ grows exponentially with k , and is roughly polynomial for $GSA - \varphi_{Gs}$ and $GSA - \varphi_{Gs+Eig}$. On the contrary, it is almost constant for $GSA - \varphi_{OPU}$ and significantly the lowest when $k \geq 7$.

To summarize, $GSA - \varphi_{OPU}$ outperforms the traditional methods both in accuracy and computation time.

4.3. Varying m, k, S_k in $GSA - \varphi_{OPU}$

From Fig. 2, we note that RW sampling, as expected, provide better results when the subgraph size k is smaller or equal to 4. Moreover, when increasing m and fixing other parameters, similarly noted with k , accuracy curves converge to the best curve that corresponds to $m \rightarrow \infty$.

4.4. Comparing $GSA - \varphi_{OPU}$ against GIN-based model

In Fig 2, We see that $GSA - \varphi_{OPU}$ with either (RW sampling, $k \geq 4$) or (uniform sampling, $k \geq 5$) performs better than the GIN-based graph convolutional model. We note that we do not report the computational time for GIN, since it is highly dependent on high-speed graphical processing units (GPUs) to do the training process.

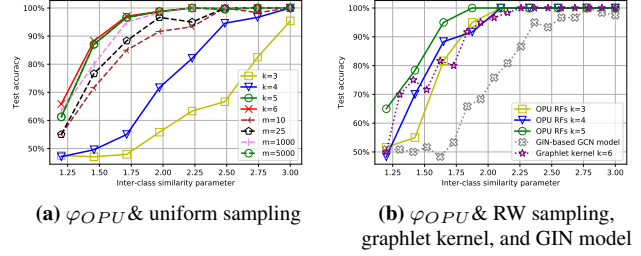


Fig. 2. Trying different values of the parameters (k, m, S_k) with $GSA - \varphi_{OPU}$. Then, benchmarking the test accuracies against the ones of $GSA - \varphi_k^{match}$ and GIN-based GCN network. We fix $s = 2000$ and $m = 5000$ while varying the problem difficulty r . The GIN-based model consists of 5 GIN layers followed by 2 fully connected layers, the dimension of hidden layers= 4.




Fig. 3. Comparing $GSA - \varphi$ against the graphlet kernel as a performance reference on real datasets with: $s=4000$, $k=7$.

4.5. $GSA - \varphi_{OPU}$ on real datasets

In Fig 3, we have the test accuracy with varying value of m . For each value of m we conduct the experiment 3 times on D&D and 4 times on Reddit-Binary dataset and take the average accuracy. For D&D, although results of the 3 experiments get more concentrated as m grows, we do not observe a steady improvement in the average accuracy. This might be accentuated by the fact that node features are ignored. However, this average is still better than the accuracy obtained by $GSA - \varphi_k^{match}$. For Reddit-Binary, the variance of experiments also decreases slightly with m . More importantly, the average accuracy is monotonically increasing, and it is larger than the accuracy of $GSA - \varphi_k^{match}$ when $m \geq 5000$.

5. CONCLUSION

We proposed a generic framework that can deploy OPUs random features in graph classification, since OPUs compute such features in $\mathcal{O}(1)$ in both input/output dimensions. Then, we showed a concentration of the random embedding around the MMD metric. Our experiments showed that our algorithm is significantly faster than the graphlet kernel with graphlet

sampling and performs better while concentrating around the MMD metric. Moreover, it outperformed a state-of-the-art graph convolutional network on graph classification.

A major point left open to be analyzed is how to use our algorithm to classify graphs with node features. One promising possibility is to use our algorithm to generate features embeddings on the graph level, and then feed these embeddings with the nodes' features to a deep neural network. On the theoretical side, the properties of the MMD metric could be further analyzed on particular models of graphs to get a concentration with higher certainty.

6. REFERENCES

- [1] Giannis Nikolentzos, Polykarpos Meladianos, and Michalis Vazirgiannis, "Matching node embeddings for graph similarity," in *Thirty-First AAAI Conference on Artificial Intelligence*, 2017.
- [2] Pinar Yanardag and SVN Vishwanathan, "Deep graph kernels," in *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2015, pp. 1365–1374.
- [3] Christopher Morris, Nils M Kriege, Franka Bause, Kristian Kersting, Petra Mutzel, and Marion Neumann, "Tudataset: A collection of benchmark datasets for learning with graphs," *arXiv preprint arXiv:2007.08663*, 2020.
- [4] Chi Thang Duong, Thanh Dat Hoang, Ha The Hien Dang, Quoc Viet Hung Nguyen, and Karl Aberer, "On node features for graph neural networks," *arXiv preprint arXiv:1911.08795*, 2019.
- [5] X Yan and J Han, "gspan: Graph-based substructure pattern mining, 2002," *Published by the IEEE Computer Society*, 2003.
- [6] Nils M Kriege, Fredrik D Johansson, and Christopher Morris, "A survey on graph kernels," *Applied Network Science*, vol. 5, no. 1, pp. 1–42, 2020.
- [7] Nino Shervashidze, SVN Vishwanathan, Tobias Petri, Kurt Mehlhorn, and Karsten Borgwardt, "Efficient graphlet kernels for large graph comparison," in *Artificial Intelligence and Statistics*, 2009, pp. 488–495.
- [8] Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka, "How powerful are graph neural networks?," *arXiv preprint arXiv:1810.00826*, 2018.
- [9] Aman Sinha and John C Duchi, "Learning kernels with random features," in *Advances in Neural Information Processing Systems*, 2016, pp. 1298–1306.
- [10] Johannes Kobler, Uwe Schöning, and Jacobo Torán, *The graph isomorphism problem: its structural complexity*, Springer Science & Business Media, 2012.
- [11] Anna Lubiw, "Some np-complete problems similar to graph isomorphism," *SIAM Journal on Computing*, vol. 10, no. 1, pp. 11–21, 1981.
- [12] Jure Leskovec and Christos Faloutsos, "Sampling from large graphs," in *Proceedings of the 12th ACM SIGKDD international conference on Knowledge discovery and data mining*, 2006, pp. 631–636.
- [13] Ali Rahimi and Benjamin Recht, "Random features for large-scale kernel machines," in *Advances in neural information processing systems*, 2008, pp. 1177–1184.
- [14] Dougal J Sutherland, Hsiao-Yu Tung, Heiko Strathmann, Soumyajit De, Aaditya Ramdas, Alex Smola, and Arthur Gretton, "Generative models and model criticism via optimized maximum mean discrepancy," *arXiv preprint arXiv:1611.04488*, 2016.
- [15] Alaa Saade, Francesco Caltagirone, Igor Carron, Laurent Daudet, Angélique Dréméau, Sylvain Gigan, and Florent Krzakala, "Random projections through multiple optical scattering: Approximating kernels at the speed of light," in *2016 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2016, pp. 6215–6219.
- [16] Thorben Funke and Till Becker, "Stochastic block models: A comparison of variants and inference methods," *PloS one*, vol. 14, no. 4, 2019.
- [17] Paul D Dobson and Andrew J Doig, "Distinguishing enzyme structures from non-enzymes without alignments," *Journal of molecular biology*, vol. 330, no. 4, pp. 771–783, 2003.

Appendices

A. PROOF OF THEOREM 1

Proof. We decompose the proof in two steps.

Step 1: infinite s , finite m . First we define the random variables $x_j = |\mathbb{E}_{F \sim S_k(\mathcal{G})} \xi_{w_j}(F) - \mathbb{E}_{F' \sim S_k(\mathcal{G}')} \xi_{w_j}(F')|^2$, which are: i/independent, ii/have expectation $MMD(\mathcal{G}, \mathcal{G}')^2$, /iii are bounded by the interval $[0, 4]$ based on our assumption $|\xi_w| \leq 1$. Thus, as a straight result of applying Hoeffding's inequality with easy manipulation: with probability $1 - \delta$

$$\left| \frac{1}{m} \sum_{j=1}^m x_j - MMD(\mathcal{G}, \mathcal{G}')^2 \right| \leq \frac{4\sqrt{\log(2/\delta)}}{\sqrt{m}} \quad (6)$$

Step 2: finite s and m . For any *fixed* set of random features $\{w_j\}_{1, \dots, m}$ and based on our previous assumptions we have: i/ φ_{RF} is in a ball of radius $M = \frac{\sqrt{m}}{\sqrt{m}} = 1$, ii/ $\mathbb{E}_{F \sim S_k(\mathcal{G})} \varphi(F) = \mathbb{E} \left(\frac{1}{s} \sum_i \varphi(F_i) \right)$. Therefore, we can directly apply the vector version of Hoeffding's inequality on the vectors $\frac{1}{s} \sum_i \varphi(F_i)$ to get that with probability $1 - \delta$:

$$\left\| \mathbb{E}_{F \sim S_k(\mathcal{G})} \varphi(F) - \frac{1}{s} \sum_i \varphi(F_i) \right\| \leq \frac{1 + \sqrt{2 \log \frac{1}{\delta}}}{\sqrt{s}} \quad (7)$$

Defining $J_{exp}(\mathcal{G}, \mathcal{G}') = \|\mathbb{E}_{F \sim S_k(\mathcal{G})} \varphi(F) - \mathbb{E}_{F' \sim S_k(\mathcal{G}')} \varphi(F')\|$ and $J_{avg}(\mathcal{G}, \mathcal{G}') = \|\frac{1}{s} \sum_i \varphi(F_i) - \frac{1}{s} \sum_i \varphi(F'_i)\|$, then using triangular inequality followed by a union bound based on (7), we have the following with probability $1 - 2\delta$,

$$|J_{exp}(\mathcal{G}, \mathcal{G}') - J_{avg}(\mathcal{G}, \mathcal{G}')| \leq \frac{2}{\sqrt{s}} \left(1 + \sqrt{2 \log \frac{1}{\delta}} \right)$$

On the other hand, $J_{exp}(\mathcal{G}, \mathcal{G}') + J_{avg}(\mathcal{G}, \mathcal{G}') \leq 4$, so with same probability:

$$|J_{exp}(\mathcal{G}, \mathcal{G}')^2 - J_{avg}(\mathcal{G}, \mathcal{G}')^2| \leq \frac{8}{\sqrt{s}} \left(1 + \sqrt{2 \log \frac{1}{\delta}} \right) \quad (8)$$

Since it is valid for any fixed set of random features, it is also valid with *joint* probability on random features and samples, by the law of total probability.

Finally, combining (6), (8) with a union bound and a triangular inequality, we have with probability $1 - 3\delta$,

$$\left| \|\varphi(\mathfrak{F}_{\mathcal{G}}) - \varphi(\mathfrak{F}_{\mathcal{G}'})\|^2 - MMD(\mathcal{G}, \mathcal{G}')^2 \right| \leq \frac{4\sqrt{\log(2/\delta)}}{\sqrt{m}} + \frac{8}{\sqrt{s}} \left(1 + \sqrt{2 \log \frac{1}{\delta}} \right)$$

which concludes the proof by taking δ as $\delta/3$. \square